

# capstone\_project

December 1, 2025

## 1 Adult Census Income Data Analysis

By Brandon Cabrera

This analysis is based on the Adult Census Income Data  
<https://www.kaggle.com/datasets/uciml/adult-census-income/data>

## 2 Adult Census Income Data Columns

The Adult Census Income data has 15 columns:

1. age: Age of the individual
2. workclass: Type of employment
3. fnlwgt: Indicates how many people the observation represents from the U.S. population
4. education: Highest level of education the individual completed
5. education.num: Numerical representation of education level
6. marital.status: Marital status of the individual
7. occupation: Occupation the individual holds
8. relationship: Relationship within the household
9. race: Race of the individual
10. sex: Biological sex
11. capital.gain: money earned from investments
12. capital.loss: money lost from investments
13. hours.per.week: Average number of hours worked per week
14. native.country: Country of origin
15. income: Indicates whether individual's income is > 50k or <= 50k

## 3 Analysis Questions

Throughout this analysis I will try to answer the following questions:

1. What was the average number of hours worked for a person who made  $> 50k$ , and what was the average number of hours worked for a person who made  $\leq 50k$ ?
2. For a person whose highest education level is only high school, who makes  $> 50k$ , if any, what is the average number of hours worked?
3. What is the average education level for a person making  $> 50k$ , and what is the average education level for a person making  $\leq 50k$ ?
4. What is the correlation between capital gain and whether or not a person makes over  $50k$ ?
5. What marital status has the most people making over  $50k$ , and what marital status has the most people making less than or equal to  $50k$ ?
6. What numerical column has the highest amount of correlation, in terms of magnitude, with whether or not a person made  $> 50k$  and whether they make  $\leq 50k$ ?
7. What workclass has the most people making  $\leq 50k$ , and what work class has the most people making  $> 50k$ ?
8. Do women or men tend to make  $> 50k$  more than the other, and do women or men tend to make  $\leq 50k$ ?
9. What is the average age of men who make  $> 50k$ , and what is the average age of women who make  $> 50k$ ?
10. Which race tends to make  $> 50k$ , more than the other races, and what race tends to make  $\leq 50k$  more than the other races?

## 4 Importing Required Libraries

Before we start loading the data we need to import important libraries

```
[2]: print("Done By Brandon Cabrera")
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
sns.set_theme()
```

Done By Brandon Cabrera

## 5 Importing Dataset

All of our data is contained within one dataset so let's load it and save it as a DataFrame

```
[3]: print("Done By Brandon Cabrera")
adult_census_data = pd.read_csv('data/adult_census_income.csv') #load the csv file into DataFrame
adult_census_data.head() # Display the first 5 rows of the dataset
```

Done By Brandon Cabrera

```
[3]:    age workclass fnlwgt      education education.num marital.status \
0     90      ?   77053      HS-grad                  9      Widowed
1     82  Private  132870      HS-grad                  9      Widowed
2     66      ?  186061  Some-college                 10      Widowed
3     54  Private  140359      7th-8th                  4      Divorced
4     41  Private  264663  Some-college                 10      Separated

          occupation relationship    race      sex capital.gain \
0           ?  Not-in-family  White  Female                  0
1  Exec-managerial  Not-in-family  White  Female                  0
2           ?      Unmarried  Black  Female                  0
3  Machine-op-inspct      Unmarried  White  Female                  0
4    Prof-specialty      Own-child  White  Female                  0

  capital.loss  hours.per.week native.country income
0        4356                40  United-States  <=50K
1        4356                18  United-States  <=50K
2        4356                40  United-States  <=50K
3        3900                40  United-States  <=50K
4        3900                40  United-States  <=50K
```

Let's get some info about the columns of the dataset including how many non-null values they have and what each respective column's data type is.

```
[4]: print("Done By Brandon Cabrera")
adult_census_data.info() #display the columns, non-null count, and data type
```

```
Done By Brandon Cabrera
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   age               32561 non-null   int64  
 1   workclass         32561 non-null   object  
 2   fnlwgt            32561 non-null   int64  
 3   education         32561 non-null   object  
 4   education.num     32561 non-null   int64  
 5   marital.status    32561 non-null   object  
 6   occupation        32561 non-null   object  
 7   relationship      32561 non-null   object  
 8   race               32561 non-null   object  
 9   sex                32561 non-null   object  
 10  capital.gain     32561 non-null   int64  
 11  capital.loss      32561 non-null   int64  
 12  hours.per.week   32561 non-null   int64  
 13  native.country    32561 non-null   object  
 14  income             32561 non-null   object
```

```
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

## 6 Cleaning and preparing dataset

Before starting to explore the dataset more in detail, let's make sure the dataset is clean, so let's check to see if there are any null values.

```
[5]: print(adult_census_data.isnull().sum()) # creates a boolean mask of the DataFrame then sum up the amount of True values(null values)
print("Done By Brandon Cabrera")
```

```
age          0
workclass    0
fnlwgt       0
education    0
education.num 0
marital.status 0
occupation   0
relationship  0
race         0
sex          0
capital.gain 0
capital.loss  0
hours.per.week 0
native.country 0
income        0
dtype: int64
Done By Brandon Cabrera
```

It appears there isn't an null values at first glace. However the dataset does contain '?' in some columns , e.g. row 0 of the workclass column contains a '?'. A question mark is a bit ambiguous so for the sake of less ambiguity let's replace the question marks for NaN. So now the data set does have null values which be dealt with later on

```
[6]: print("Done By Brandon Cabrera")
adult_census_data.replace('?', np.nan, inplace=True) # every '?' value will be replaced for NaN
adult_census_data.head(10) # display first 10 rows
```

Done By Brandon Cabrera

```
[6]:   age   workclass  fnlwgt   education  education.num marital.status \
0    90      NaN     77053    HS-grad           9      Widowed
1    82    Private   132870    HS-grad           9      Widowed
2    66      NaN    186061  Some-college        10      Widowed
3    54    Private   140359    7th-8th          4      Divorced
4    41    Private   264663  Some-college        10  Separated
5    34    Private   216864    HS-grad           9      Divorced
```

```

6   38      Private  150601          10th           6       Separated
7   74      State-gov  88638        Doctorate        16    Never-married
8   68      Federal-gov 422013        HS-grad         9      Divorced
9   41      Private   70037  Some-college      10    Never-married

          occupation      relationship     race      sex  capital.gain \
0            NaN  Not-in-family  White  Female           0
1  Exec-managerial  Not-in-family  White  Female           0
2            NaN      Unmarried  Black  Female           0
3  Machine-op-inspct      Unmarried  White  Female           0
4  Prof-specialty      Own-child  White  Female           0
5  Other-service      Unmarried  White  Female           0
6  Adm-clerical      Unmarried  White   Male            0
7  Prof-specialty  Other-relative  White  Female           0
8  Prof-specialty  Not-in-family  White  Female           0
9  Craft-repair      Unmarried  White   Male            0

  capital.loss  hours.per.week native.country income
0        4356                  40 United-States <=50K
1        4356                  18 United-States <=50K
2        4356                  40 United-States <=50K
3        3900                  40 United-States <=50K
4        3900                  40 United-States <=50K
5        3770                  45 United-States <=50K
6        3770                  40 United-States <=50K
7        3683                  20 United-States >50K
8        3683                  40 United-States <=50K
9        3004                  60           NaN >50K

```

```
[7]: print("Done By Brandon Cabrera")
adult_census_data.tail(10) # display the last 10 rows
```

Done By Brandon Cabrera

```
[7]: age      workclass  fnlwgt      education  education.num \
32551  43  Self-emp-not-inc  27242  Some-college           10
32552  32      Private   34066          10th            6
32553  43      Private   84661  Assoc-voc            11
32554  32      Private  116138    Masters            14
32555  53      Private  321865    Masters            14
32556  22      Private  310152  Some-college           10
32557  27      Private  257302  Assoc-acdm            12
32558  40      Private  154374    HS-grad            9
32559  58      Private  151910    HS-grad            9
32560  22      Private  201490    HS-grad            9

  marital.status      occupation      relationship \
32551 Married-civ-spouse  Craft-repair        Husband
```

32552	Married-civ-spouse	Handlers-cleaners	Husband			
32553	Married-civ-spouse	Sales	Husband			
32554	Never-married	Tech-support	Not-in-family			
32555	Married-civ-spouse	Exec-managerial	Husband			
32556	Never-married	Protective-serv	Not-in-family			
32557	Married-civ-spouse	Tech-support	Wife			
32558	Married-civ-spouse	Machine-op-inspct	Husband			
32559	Widowed	Adm-clerical	Unmarried			
32560	Never-married	Adm-clerical	Own-child			
	race	sex	capital.gain	capital.loss	hours.per.week	\
32551	White	Male	0	0	50	
32552	Amer-Indian-Eskimo	Male	0	0	40	
32553	White	Male	0	0	45	
32554	Asian-Pac-Islander	Male	0	0	11	
32555	White	Male	0	0	40	
32556	White	Male	0	0	40	
32557	White	Female	0	0	38	
32558	White	Male	0	0	40	
32559	White	Female	0	0	40	
32560	White	Male	0	0	20	
	native.country	income				
32551	United-States	<=50K				
32552	United-States	<=50K				
32553	United-States	<=50K				
32554	Taiwan	<=50K				
32555	United-States	>50K				
32556	United-States	<=50K				
32557	United-States	<=50K				
32558	United-States	>50K				
32559	United-States	<=50K				
32560	United-States	<=50K				

Let's see how many null values we have now after replacing the '?' values.

```
[8]: print("Done By Brandon Cabrera")
print(adult_census_data.isnull().sum()) # creates a boolean mask of the DataFrame then sum up the amount of True values(null values)
```

Done By Brandon Cabrera

age	0
workclass	1836
fnlwgt	0
education	0
education.num	0
marital.status	0
occupation	1843

```
relationship      0
race              0
sex               0
capital.gain     0
capital.loss     0
hours.per.week   0
native.country    583
income            0
dtype: int64
```

Notice that the missing values are only for qualitative features and that all of the numeric features don't have any missing values. This sheds some light on why originally there were '?' values for these columns instead of NaN, because it's a qualitative column, and NaN isn't the most appropriate placeholder. For a better understanding of how many null values, let's have a look at the percentage of null values there is.

```
[9]: print("Done By Brandon Cabrera")
print(((adult_census_data.isnull().sum() / len(adult_census_data)) * 100).
      round(2)) # divide null values by total number of values in column and round
      to two decimals
```

```
Done By Brandon Cabrera
age             0.00
workclass       5.64
fnlwgt          0.00
education       0.00
education.num   0.00
marital.status  0.00
occupation      5.66
relationship    0.00
race            0.00
sex             0.00
capital.gain   0.00
capital.loss   0.00
hours.per.week 0.00
native.country  1.79
income          0.00
dtype: float64
```

The highest percentage is in the occupation column, meaning 5.66 % of the values in the feature are null values. The percentage isn't that high, so we can get rid of the rows with null values without worrying about losing too much data.

```
[10]: print("Done By Brandon Cabrera")
adult_census_data.dropna(inplace = True)
adult_census_data.count()
```

```
Done By Brandon Cabrera
```

```
[10]: age          30162
       workclass    30162
       fnlwgt       30162
       education     30162
       education.num 30162
       marital.status 30162
       occupation     30162
       relationship    30162
       race           30162
       sex             30162
       capital.gain   30162
       capital.loss    30162
       hours.per.week 30162
       native.country  30162
       income          30162
       dtype: int64
```

The null values in the dataset have been removed. I now want to add two binary columns using the ‘income’ column. One column will represent if the individual makes over 50k, and the other column will represent if the individual makes less than or equal to 50k.

```
[11]: print("Done By Brandon Cabrera")
def add_income_binary_columns(df:pd.DataFrame) -> pd.DataFrame:
    """Adds two binary columns based on the income column to DataFrame passed as
    an argument. One column represents if the individual makes
    > $50k and the other column represents whether the individual makes <= 50k.

    Args:
        df (DataFrame): DataFrame to add columns to

    Returns:
        DataFrame: Returns a DataFrame with the original columns and two new
        binary columns
    """
    df['binary_income_over_$50k'] = (adult_census_data['income'] == '>50K').
    astype(int) # 1 for over 50k and 0 for <= 50k
    df['binary_income_equal_under_$50k'] = (adult_census_data['income'] ==
    '<=50K').astype(int) # 1 for <= 50k and 0 for > 50k
    return df
```

Done By Brandon Cabrera

Let’s add the columns and then confirm that the two new columns were added.

```
[12]: print("Done By Brandon Cabrera")
adult_census_data = add_income_binary_columns(adult_census_data)
```

```
adult_census_data[['binary_income_over_50k',  
                   'binary_income_equal_under_50k']].head(10) # display first 10 rows
```

Done By Brandon Cabrera

```
[12]:    binary_income_over_50k  binary_income_equal_under_50k  
1                  0                  1  
3                  0                  1  
4                  0                  1  
5                  0                  1  
6                  0                  1  
7                  1                  0  
8                  0                  1  
10                 1                  0  
11                 1                  0  
12                 1                  0
```

We are going to drop 'fnlwgt' column since it doesn't provide us any useful information to help answer our questions

```
[13]: print("Done By Brandon Cabrera")  
adult_census_data.drop('fnlwgt', inplace = True, axis = 1)
```

Done By Brandon Cabrera

```
[14]: print("Done By Brandon Cabrera")  
adult_census_data.head()
```

Done By Brandon Cabrera

```
[14]:    age workclass      education  education.num marital.status  \  
1     82   Private       HS-grad          9      Widowed  
3     54   Private       7th-8th          4      Divorced  
4     41   Private     Some-college        10     Separated  
5     34   Private       HS-grad          9      Divorced  
6     38   Private       10th            6     Separated  
  
          occupation  relationship    race      sex  capital.gain  \  
1 Exec-managerial Not-in-family  White Female          0  
3 Machine-op-inspct      Unmarried  White Female          0  
4 Prof-specialty       Own-child  White Female          0  
5 Other-service        Unmarried  White Female          0  
6 Adm-clerical         Unmarried  White  Male           0  
  
    capital.loss  hours.per.week native.country income  \  
1          4356              18 United-States <=50K  
3          3900              40 United-States <=50K  
4          3900              40 United-States <=50K  
5          3770              45 United-States <=50K
```

```

6           3770          40  United-States  <=50K

  binary_income_over_>$50k  binary_income_equal_under_>$50k
1                  0                      1
3                  0                      1
4                  0                      1
5                  0                      1
6                  0                      1

```

## 7 Exploratory Data Analysis(EDA)

The exploration can be split up into two categories: numerical columns and qualitative columns. Our goal here is to get some graphs that help answer our analysis questions.

### 7.1 Exploring numerical columns

We're going to do some boxplots for our numerical columns as the y-values for our boxplots, and then for our x-values we'll do one of the binary income columns. It doesn't matter which of the two columns we choose since we'll still get the same information, the graphs would just be mirrored versions of each other.

```
[15]: print("Done by Brandon Cabrera")
adult_census_data.select_dtypes(include='number').columns # show all of our
↳ numeric columns
```

Done by Brandon Cabrera

```
[15]: Index(['age', 'education.num', 'capital.gain', 'capital.loss',
       'hours.per.week', 'binary_income_over_>$50k',
       'binary_income_equal_under_>$50k'],
      dtype='object')
```

```
[16]: print("Done By Brandon Cabrera")
def plot_boxplot(x:str, y:str) -> pd.DataFrame:
    """Plots a boxplot and returns a DataFrame containing descriptive stats shown
    by the boxplot

    Args:
        x (String): A column name from adult_census_data
        y (String): A column name from adult_census_data

    Returns:
        DataFrame: Contains descriptive stats of what's shown on the boxplot
    """
    fig = plt.figure(figsize=(12,12))
    fig = sns.boxplot(data = adult_census_data, x = adult_census_data[x], y = adult_census_data[y], hue = adult_census_data['income'],)
    plt.show()
```

```


group_0 = adult_census_data[adult_census_data['binary_income_over_$50k'] == 0][y]
group_1 = adult_census_data[adult_census_data['binary_income_over_$50k'] == 1][y]
group_combined = pd.DataFrame({'Income <= 50k(0)': group_0.describe(), 'Income > 50k(1)': group_1.describe()})
return group_combined

```

Done By Brandon Cabrera

```
[17]: print("Done By Brandon Cabrera")
hours_per_week_info = plot_boxplot(x = 'binary_income_over_$50k', y = 'hours.per.week')
hours_per_week_info
```

Done By Brandon Cabrera




```
[17]:      Income <= 50k(0)  Income > 50k(1)
count      22654.000000    7508.000000
mean       39.348592     45.706580
std        11.950774     10.736987
min        1.000000      1.000000
25%       38.000000     40.000000
50%       40.000000     40.000000
75%       40.000000     50.000000
max       99.000000     99.000000
```

From the above boxplot, it's clear that people who made  $> 50k$  had a higher average number of hours worked per week, with 45.7 hours worked. Their standard deviation are similar, with a difference of around 1.22 hours worked. The quartiles for the two groups are close in value as well. This gives us enough information to answer the analysis question #1.

```
[18]: print("Done By Brandon Cabrera")
capital_loss_info = plot_boxplot(x = 'binary_income_over_$50k', y = 'capital.
    ↴loss')
capital_loss_info
```

Done By Brandon Cabrera




```
[18]:      Income <= 50k(0)  Income > 50k(1)
count      22654.000000    7508.000000
mean       53.448000     193.750666
std        310.270263    592.825590
min        0.000000     0.000000
25%        0.000000     0.000000
50%        0.000000     0.000000
75%        0.000000     0.000000
max       4356.000000    3683.000000
```

The quartiles of the boxplot are hard to see since it's been squished to the bottom of the graph. From the descriptive stats, people who make over 50k have a mean of 193.75 capital loss, so they actually lose more from investments than people who make under 50k, with a mean of 53.44 capital loss. The max value for people who make  $\leq 50k$  actually is higher than those who do make  $> 50k$ .

50k with values of 4356.0 and 3683.0, respectively.

```
[43]: print("Done By Brandon Cabrera")
capital_gain_info = plot_boxplot(x = 'binary_income_over_$50k', y = 'capital.
           gain')
capital_gain_info
```

Done By Brandon Cabrera




```
[43]:      Income <= 50k(0)  Income > 50k(1)
count      22654.000000    7508.000000
mean       148.893838    3937.679808
std        936.392280   14386.060019
min        0.000000     0.000000
25%       0.000000     0.000000
```

50%	0.000000	0.000000
75%	0.000000	0.000000
max	41310.000000	99999.000000

Once again, the quartiles for the boxplot have been squished to the bottom of the graph. The descriptive stats show that for the people who make  $> 50k$ , they have a mean of 3937.67 capital gain, while the people who make  $\leq 50k$  have a mean of 148.89 capital gain. That is a sizeable difference between the two. The max value of capital gain is 41310.0 for people who make  $< 50k$ , and then for people who make  $> 50k$  it's 99999.0

```
[20]: print("Done By Brandon Cabrera")
age_group_info = plot_boxplot(x = 'binary_income_over_$50k', y='age')
age_group_info
```

Done By Brandon Cabrera




```
[20]: Income <= 50k(0)  Income > 50k(1)
      count    22654.000000    7508.000000
      mean     36.608060     43.959110
      std      13.464631     10.269633
      min      17.000000     19.000000
      25%     26.000000     36.000000
      50%     34.000000     43.000000
      75%     45.000000     51.000000
      max     90.000000     90.000000
```

From the above figure and the statistics, the mean age of people who make > 50k is higher than the mean age of people who make <= 50k. With values of 43.95 years old and 36.60 years old, respectively.

```
[21]: print("Done By Brandon Cabrera")
education_num_info = plot_boxplot(x = 'binary_income_over_$50k', y = 'education.
˓→num')
education_num_info
```

Done By Brandon Cabrera



```
[21]:      Income <= 50k(0)  Income > 50k(1)
count      22654.000000    7508.000000
mean       9.629116     11.606420
std        2.413596     2.368423
min        1.000000     2.000000
25%        9.000000    10.000000
50%        9.000000    12.000000
75%        10.000000   13.000000
max       16.000000    16.000000
```

The above information tells us that the average education number for people who make > 50k is 12, if we round up to the nearest number, which is equivalent to Assoc-acdm in terms of education level. Then, for the people who make <= 50k, rounding up to the nearest whole number, the average education number is 10, meaning they have some-college. This answers analysis question

#3, the average education level for both groups is now known.

### 7.1.1 Numerical Columns Correlations

Next let's see how correlated the numerical columns are to each binary column.

```
[22]: print("Done By Brandon Cabrera")
adult_census_data.corr(numeric_only= True)[['binary_income_over_50k',\n    'binary_income_equal_under_50k']]
```

Done By Brandon Cabrera

```
[22]:                                binary_income_over_50k \
age                               0.241998
education.num                      0.335286
capital.gain                        0.221196
capital.loss                         0.150053
hours.per.week                      0.229480
binary_income_over_50k                1.000000
binary_income_equal_under_50k        -1.000000
```

```
                                binary_income_equal_under_50k
age                           -0.241998
education.num                  -0.335286
capital.gain                   -0.221196
capital.loss                    -0.150053
hours.per.week                 -0.229480
binary_income_over_50k            -1.000000
binary_income_equal_under_50k      1.000000
```

From the correlation matrix analysis, question #4 and analysis question #6 can be answered. The correlation between capital gain and a person making over > 50k is 0.221, indicating a weak positive relationship. This makes sense since not everyone who is making money from investments might be making a profit compared to their losses, or the money made isn't much. The column with the highest amount of correlation, in terms of magnitude, for both binary columns is education.num with a value of 0.335. This is a weak relationship, indicating that as a person's education level becomes higher, they are more likely to make >50k. If we look at it from the perspective of making <= 50k, then it indicates that as a person's education level becomes higher than they are less likely to make <= 50k. Capital loss has the least amount of correlation with both of the income groups.

## 7.2 Exploring qualitative columns

For visualizing the qualitative columns let's use bar chart to help answer the analysis questions. Let's get a list of the qualitative columns.

```
[23]: print("Done by Brandon Cabrera")
adult_census_data.select_dtypes(include = 'object').columns
```

Done by Brandon Cabrera

```
[23]: Index(['workclass', 'education', 'marital.status', 'occupation',
       'relationship', 'race', 'sex', 'native.country', 'income'],
       dtype='object')
```

```
[24]: print("Done by Brandon Cabrera")
def plot_countplot(col: str, figsize = (15,15)) -> pd.DataFrame:
    """Creates a countplot and returns a frequency table that represents the
    countplot


    Args:
        col (str): A column name from adult_census_data

    Returns:
        pd.DataFrame: A Frequency table that represents the countplot that is made
    """
    fig = plt.figure(figsize = figsize)
    fig = sns.countplot(x=adult_census_data[col], hue =
    adult_census_data['income'])
    plt.show()
    return pd.crosstab(index= adult_census_data[col], columns =
    adult_census_data['binary_income_over_$50k']) #creates a frequency table
```

Done by Brandon Cabrera

```
[25]: print("Done by Brandon Cabrera")
workclass_countplot_table = plot_countplot('workclass')
workclass_countplot_table.sort_values(by=0, ascending = False) # sort by column
# 0 and go from largest to smallest
```

Done by Brandon Cabrera




```
[25]: binary_income_over_50k      0      1
       workclass
       Private      17410  4876
       Self-emp-not-inc  1785  714
       Local-gov     1458  609
       State-gov      935  344
       Federal-gov    578  365
       Self-emp-inc    474  600
       Without-pay     14   0
```

The above figure shows that the private workclass has the most amount of people making  $> 50k$ , as well as the most amount of people making  $\leq 50k$ . There are 4876 people who make  $> 50k$  and who are a part of the private workclass. This answers analysis question #7. A trend for the data is that the  $\leq 50k$  income group has more people in it than the  $> 50k$  income group for

every workclass, except for the self-emp-inc workclass.

```
[26]: print("Done by Brandon Cabrera")
education_countplot_table = plot_countplot('education')
education_countplot_table.sort_values(by=0, ascending = False) # sort by column
    ↪ 0 and go from largest to smallest
```

Done by Brandon Cabrera




```
[26]: binary_income_over_50k      0      1
      education
      HS-grad          8223  1617
      Some-college     5342  1336
      Bachelors        2918  2126
```

11th	989	59
Assoc-voc	963	344
10th	761	59
Assoc-acdm	752	256
Masters	709	918
7th-8th	522	35
9th	430	25
12th	348	29
5th-6th	276	12
1st-4th	145	6
Prof-school	136	406
Doctorate	95	280
Preschool	45	0

It's clear that education level with the most amount of people who are making  $\leq 50k$  is the HS-grad education level, from the above figure. We can also see that for post-graduate education levels, there is more people who make over  $>50k$  than there are people who make  $\leq 50k$ . We can see the opposite effect for people whose highest education level is a Bachelor's degree or lower.

```
[27]: print("Done by Brandon Cabrera")
marital_status_countplot_table = plot_countplot('marital.status')
marital_status_countplot_table.sort_values(by=0, ascending = False) # sort by
#column 0 and go from largest to smallest
```

Done by Brandon Cabrera




```
[27]: binary_income_over_50k      0      1
      marital.status
      Never-married      9256   470
      Married-civ-spouse  7666  6399
      Divorced          3762   452
      Separated          873    66
      Widowed            747    80
      Married-spouse-absent 339    31
      Married-AF-spouse   11    10
```

From the above figure, analysis question #5 can be answered. The marital status with the most amount of people making > 50k is married-civ-spouse. The marital status with the most amount of people making <= 50k is Never-married.

```
[28]: print("Done by Brandon Cabrera")
occupation_countplot_table = plot_countplot('occupation', figsize= (24,15))
occupation_countplot_table.sort_values(by=0, ascending = False) # sort by
    ↪column 0 and go from largest to smallest
```

Done by Brandon Cabrera




```
[28]: binary_income_over_>50k      0      1
occupation
Adm-clerical          3223  498
Craft-repair            3122  908
Other-service           3080  132
Sales                   2614  970
Prof-specialty          2227  1811
Exec-managerial         2055  1937
Machine-op-inspct       1721  245
Handlers-cleaners       1267  83
Transport-moving        1253  319
Farming-fishing         874   115
Tech-support             634   278
Protective-serv         434   210
Priv-house-serv         142   1
Armed-Forces              8   1
```

From the above figure, it's evident that most of the occupations have a considerable amount of more

people making  $\leq 50k$  than people making  $\geq 50k$ . Except for the occupations, Prof-specialty and Exec-managerial, where the difference is a lot smaller.

```
[29]: print("Done by Brandon Cabrera")
relationship_countplot_table = plot_countplot(col = 'relationship')
relationship_countplot_table.sort_values(by=0, ascending = False) # sort by
    ↵column 0 and go from largest to smallest
```

Done by Brandon Cabrera




```
[29]: binary_income_over_>$50k      0      1
      relationship
      Not-in-family      6903     823
      Husband            6784    5679
```

Own-child	4402	64
Unmarried	2999	213
Other-relative	854	35
Wife	712	694

The relationship columns represent the individual's relationship inside the household, so it makes sense that the husband and wife relationship is the one with the most amount of people making > 50k as opposed to the other relationship statuses.

```
[30]: print("Done by Brandon Cabrera")
race_countplot_table = plot_countplot(col = 'race')
race_countplot_table.sort_values(by=0, ascending = False) # sort by column 0
    ↵and go from largest to smallest
```

Done by Brandon Cabrera




```
[30]: binary_income_over_50k      0      1
race
White           19094  6839
Black            2451   366
Asian-Pac-Islander    647   248
Amer-Indian-Eskimo    252    34
Other             210    21
```

Analysis question #10 can now be answered. It is clear the race with the most amount of people making > 50k is the white race and it's also the race with the most amount of people making <= 50k. It's worth noting that I am considering only the frequency and not the percentages, which would be a fairer assessment of which race tends to make more.

```
[31]: print("Done by Brandon Cabrera")
sex_countplot_table = plot_countplot(col='sex')
sex_countplot_table.sort_values(by=0, ascending = False) # sort by column 0 and
# go from largest to smallest
```

Done by Brandon Cabrera




```
[31]: binary_income_over_50k      0      1
      sex
      Male          13984  6396
      Female         8670  1112
```

The above information shows that the sex with the most amount of people who make > 50k and the most of people who make  $\leq 50k$  is the male sex. This answers analysis question #8. In both groups, there is a considerable amount more people making  $\leq 50k$ .

```
[32]: print("Done by Brandon Cabrera")
native_country_countplot_table = plot_countplot(col = 'native.country')
native_country_countplot_table.sort_values(by=0, ascending = False) # sort by
# column 0 and go from largest to smallest
```

Done by Brandon Cabrera



```
[32]: binary_income_over_50k      0      1
native.country
United-States      20509  6995
Mexico            577   33
Philippines       128   60
Puerto-Rico       97    12
El-Salvador       91    9
Germany          84    44
Canada            71    36
Jamaica           70    10
Cuba              67    25
Dominican-Republic 65    2
Guatemala         60    3
India              60   40
Vietnam            59    5
```

South	57	14
England	56	30
Columbia	54	2
China	48	20
Poland	45	11
Italy	44	24
Haiti	38	4
Japan	36	23
Nicaragua	31	2
Portugal	30	4
Peru	28	2
Iran	24	18
Taiwan	23	19
Ecuador	23	4
Greece	21	8
Ireland	19	5
Trinidad&Tobago	16	2
Laos	15	2
France	15	12
Outlying-US(Guam-USVI-etc)	14	0
Thailand	14	3
Hong	13	6
Cambodia	11	7
Honduras	11	1
Hungary	10	3
Yugoslavia	10	6
Scotland	9	2
Holland-Netherlands	1	0

There are too many countries in the dataset to properly show on the countplot, but the frequency table shows that the United States has the most amount of people in both income categories. Mexico has the second most amount of people making  $\leq 50k$ , followed by the Philippines.

### 7.2.1 Pivot Tables for Qualitative Columns

Now I am going to make some pivot tables for the qualitative columns and for each qualitative column we will separate them into the two income groups and have them summarize the numeric columns by calculating the mean of each numeric column. Although I am going to make one for each qualitative column I'm most interested into the ones that help answer the remaining analysis questions, #3 and #9.

```
[33]: print("Done by Brandon Cabrera")
workclass_pivot_table = pd.pivot_table(data= adult_census_data, index = ['workclass', 'income'], values= ['age', 'education.num', 'capital.gain', 'capital.loss', 'hours.per.week'])
workclass_pivot_table
```

Done by Brandon Cabrera

[33] :

		age	capital.gain	capital.loss	education.num	\
workclass	income					
Federal-gov	<=50K	40.605536	176.503460	93.015571	10.468858	
	>50K	45.701370	1870.849315	141.369863	11.706849	
Local-gov	<=50K	40.714678	166.970508	80.163237	10.569959	
	>50K	44.205255	2414.738916	182.528736	12.154351	
Private	<=50K	35.106720	138.144170	49.671051	9.438369	
	>50K	42.820139	3528.184988	186.597621	11.416120	
Self-emp-inc	<=50K	43.130802	182.476793	56.662447	10.251055	
	>50K	48.316667	8467.080000	230.010000	11.891667	
Self-emp-not-inc	<=50K	44.451541	223.535014	61.985434	9.680112	
	>50K	46.411765	6137.133053	249.539216	11.539216	
State-gov	<=50K	37.240642	139.203209	40.529412	10.790374	
	>50K	45.127907	2165.909884	191.549419	12.869186	
Without-pay	<=50K	47.785714	487.857143	0.000000	9.071429	

		hours.per.week
workclass	income	
Federal-gov	<=50K	39.982699
	>50K	43.334247
Local-gov	<=50K	39.748971
	>50K	44.003284
Private	<=50K	38.782596
	>50K	45.493437
Self-emp-inc	<=50K	46.966245
	>50K	50.253333
Self-emp-not-inc	<=50K	43.507003
	>50K	46.745098
State-gov	<=50K	37.170053
	>50K	44.174419
Without-pay	<=50K	32.714286

Through the above pivot table, there is a trend that the average age for the > 50k group for every workclass is higher than the average age for the <= 50k group of the same workclass. We can also see this same trend with all the other numeric columns; people in the > 50k group have more capital gain, capital loss, higher education levels, and more hours worked per week than their counterparts in the <= 50k group.

```
[34]: print("Done by Brandon Cabrera")
education_pivot_table = pd.pivot_table(data= adult_census_data, index = [
    'education', 'income'], values= ['age', 'education.num', 'capital.gain',
    'capital.loss', 'hours.per.week'])
education_pivot_table
```

Done by Brandon Cabrera

```
[34]:                                     age  capital.gain  capital.loss  education.num \
education      income
```

10th	<=50K	36.980289	97.618922	42.540079	6.0
	>50K	49.728814	4243.423729	316.728814	6.0
11th	<=50K	31.731041	80.803842	43.089990	7.0
	>50K	42.966102	2503.406780	210.915254	7.0
12th	<=50K	30.899425	89.183908	30.321839	8.0
	>50K	45.379310	2381.310345	63.724138	8.0
1st-4th	<=50K	44.317241	65.634483	55.993103	2.0
	>50K	52.000000	1281.333333	0.000000	2.0
5th-6th	<=50K	41.554348	105.786232	67.797101	3.0
	>50K	43.833333	1647.833333	157.250000	3.0
7th-8th	<=50K	47.260536	185.296935	67.026820	4.0
	>50K	53.171429	1130.714286	54.342857	4.0
9th	<=50K	39.793023	130.300000	21.179070	5.0
	>50K	49.080000	4207.080000	163.160000	5.0
Assoc-acdm	<=50K	35.902926	114.837766	65.966755	12.0
	>50K	41.351562	1847.667969	177.835938	12.0
Assoc-voc	<=50K	36.741433	194.452752	46.686397	11.0
	>50K	42.459302	2257.125000	136.020349	11.0
Bachelors	<=50K	35.893763	165.098355	59.048321	13.0
	>50K	42.412982	3889.806209	198.828786	13.0
Doctorate	<=50K	44.357895	248.336842	51.884211	16.0
	>50K	48.071429	6654.185714	337.610714	16.0
HS-grad	<=50K	37.451417	153.723945	53.116502	9.0
	>50K	44.690167	2804.876314	160.666048	9.0
Masters	<=50K	41.895628	293.393512	83.358251	14.0
	>50K	45.164488	4300.569717	240.202614	14.0
Preschool	<=50K	41.288889	1018.177778	75.355556	1.0
Prof-school	<=50K	42.882353	156.375000	95.808824	15.0
	>50K	44.706897	14274.081281	280.923645	15.0
Some-college	<=50K	34.196181	126.955447	50.589105	10.0
	>50K	43.889222	2442.908683	157.123503	10.0

hours.per.week

education	income	
10th	<=50K	36.988173
	>50K	43.610169
11th	<=50K	33.555106
	>50K	44.898305
12th	<=50K	34.939655
	>50K	44.793103
1st-4th	<=50K	37.944828
	>50K	48.833333
5th-6th	<=50K	38.521739
	>50K	45.166667
7th-8th	<=50K	39.609195
	>50K	47.914286
9th	<=50K	38.413953

	>50K	44.840000
Assoc-acdm	<=50K	39.953457
	>50K	44.800781
Assoc-voc	<=50K	41.298027
	>50K	43.790698
Bachelors	<=50K	40.920493
	>50K	45.731891
Doctorate	<=50K	46.694737
	>50K	48.217857
HS-grad	<=50K	40.222303
	>50K	45.210884
Masters	<=50K	41.703808
	>50K	46.200436
Preschool	<=50K	36.866667
Prof-school	<=50K	43.816176
	>50K	49.352217
Some-college	<=50K	37.999064
	>50K	45.056886

The above pivot table helps to answer the analysis question #2. For a person whose highest education level is HS-grad, the average age for people of that education level who make > 50k is 45 years old, if we round up to the nearest whole number. Another pattern in this pivot table is that for the higher education levels past Some-college, the difference in the average age between the income groups is a lot smaller. The Doctorate level has a mean age of 44 years old for the <= 50k group and a mean age of 48 years old for the > 50k group. The Some-college group has a mean age of 34 years old for it's <= 50k group and a mean age of 44 years old for it's > 50k group

```
[35]: print("Done by Brandon Cabrera")
marital_status_pivot_table = pd.pivot_table(data= adult_census_data, index = ['marital.status', 'income'], values= ['age', 'education.num', 'capital.gain', 'capital.loss', 'hours.per.week'])
marital_status_pivot_table
```

Done by Brandon Cabrera

```
[35]:
```

marital.status	income	age	capital.gain	capital.loss	\
Divorced	<=50K	42.566986	142.161350	57.413610	
	>50K	45.435841	5838.396018	141.207965	
Married-AF-spouse	<=50K	30.090909	0.000000	0.000000	
	>50K	31.300000	729.800000	0.000000	
Married-civ-spouse	<=50K	41.636968	219.403861	61.375815	
	>50K	44.127989	3605.772152	198.067042	
Married-spouse-absent	<=50K	39.631268	113.784661	42.628319	
	>50K	47.645161	6600.580645	130.193548	
Never-married	<=50K	27.974395	98.092589	45.740169	
	>50K	38.053191	6011.123404	179.006383	
Separated	<=50K	39.174112	117.262314	48.349370	

	>50K	42.348485	6614.727273	232.621212
Widowed	<=50K	57.693440	143.764391	59.281124
	>50K	58.287500	4726.162500	248.762500
education.num hours.per.week				
marital.status	income			
Divorced	<=50K	9.852472	40.808612	
	>50K	11.918142	47.460177	
Married-AF-spouse	<=50K	9.363636	45.727273	
	>50K	11.000000	42.600000	
Married-civ-spouse	<=50K	9.360031	42.326768	
	>50K	11.517425	45.558681	
Married-spouse-absent	<=50K	8.932153	39.368732	
	>50K	12.032258	45.225806	
Never-married	<=50K	9.881698	36.756590	
	>50K	12.534043	46.674468	
Separated	<=50K	9.174112	39.217640	
	>50K	12.166667	46.212121	
Widowed	<=50K	8.987952	33.599732	
	>50K	10.962500	42.100000	

A trend that's seen in the above pivot table is that for the marital status groups that have at one point been married or are still currently married differ by a small amount of years in their mean age for the income groups, except for the Married-spouse-absent. The <= 50k group of the Married-spouse-absent marital status has a mean age of 40 years old, whereas the > 50k group has a mean age of 48 years old

```
[44]: print("Done by Brandon Cabrera")
occupation_pivot_table = pd.pivot_table(data= adult_census_data, index = ['occupation', 'income'], values= ['age', 'education.num', 'capital.gain', 'capital.loss', 'hours.per.week'])
occupation_pivot_table
```

Done by Brandon Cabrera

```
[44]:                                     age  capital.gain  capital.loss \
occupation      income
Adm-clerical    <=50K   36.031958   138.317096   50.753025
                  >50K   43.299197   2819.082329  119.753012
Armed-Forces     <=50K   28.250000   0.000000   0.000000
                  >50K   46.000000   0.000000  1887.000000
Craft-repair     <=50K   37.592569   148.671685   67.058616
                  >50K   43.735683   2407.403084  153.394273
Exec-managerial  <=50K   39.628224   184.769830   59.030170
                  >50K   44.893650   4307.764068  224.303562
Farming-fishing  <=50K   40.487414   266.659039   43.516018
                  >50K   47.060870   3070.678261  214.469565
Handlers-cleaners <=50K   31.385951   112.342541   36.377269
```

	>50K	43.240964	2483.746988	163.939759
Machine-op-inspct	<=50K	37.038350	155.248693	45.747240
	>50K	42.297959	1507.987755	148.293878
Other-service	<=50K	34.628896	78.926948	32.992857
	>50K	41.371212	2579.871212	127.439394
Priv-house-serv	<=50K	42.000000	115.929577	22.507042
	>50K	47.000000	25236.000000	0.000000
Prof-specialty	<=50K	37.911540	193.944320	66.916031
	>50K	43.600221	5821.582551	221.925456
Protective-serv	<=50K	37.688940	247.887097	46.822581
	>50K	41.480952	1676.061905	145.323810
Sales	<=50K	34.802984	135.007651	62.429610
	>50K	44.358763	4471.389691	198.073196
Tech-support	<=50K	34.399054	194.793375	54.121451
	>50K	43.140288	1747.528777	199.089928
Transport-moving	<=50K	39.166800	131.960096	65.486832
	>50K	44.517241	1936.366771	150.689655

occupation	income	education.num	hours.per.week
Adm-clerical	<=50K	10.013962	37.054918
	>50K	10.716867	40.839357
Armed-Forces	<=50K	9.625000	40.750000
	>50K	14.000000	40.000000
Craft-repair	<=50K	8.914478	41.612108
	>50K	9.817181	44.656388
Exec-managerial	<=50K	10.825304	42.760097
	>50K	12.089830	47.308209
Farming-fishing	<=50K	8.431350	46.041190
	>50K	10.008696	54.208696
Handlers-cleaners	<=50K	8.444357	37.581689
	>50K	9.277108	42.349398
Machine-op-inspct	<=50K	8.350959	40.336432
	>50K	9.351020	43.310204
Other-service	<=50K	8.720455	34.250000
	>50K	10.136364	42.901515
Priv-house-serv	<=50K	7.183099	32.781690
	>50K	13.000000	35.000000
Prof-specialty	<=50K	12.365963	40.071846
	>50K	13.536720	45.207068
Protective-serv	<=50K	9.845622	41.525346
	>50K	10.900000	45.576190
Sales	<=50K	9.902448	38.276970
	>50K	11.364948	47.463918
Tech-support	<=50K	10.839117	38.597792
	>50K	11.258993	41.471223
Transport-moving	<=50K	8.636073	43.619314

>50K	9.206897	48.699060
------	----------	-----------

An interesting thing to note from the above pivot table is that none of the income groups for Armed-Forces have any capital gain and the amount of hours per week are almost the same.

```
[37]: print("Done by Brandon Cabrera")
relationship_pivot_table = pd.pivot_table(data= adult_census_data, index =[['relationship', 'income'], values= ['age', 'education.num', 'capital.gain','capital.loss', 'hours.per.week']])
relationship_pivot_table
```

Done by Brandon Cabrera

```
[37]:
```

		age	capital.gain	capital.loss	education.num
relationship	income				
	Husband	<=50K	42.131191	206.296285	62.477889
	>50K	44.561719	3691.779010	200.874274	11.527558
Not-in-family	<=50K	37.327973	138.764305	65.160220	10.144865
	>50K	42.701094	5940.450790	164.340219	12.262454
Other-relative	<=50K	32.853630	129.375878	40.624122	8.679157
	>50K	42.000000	2804.428571	398.428571	10.800000
Own-child	<=50K	25.028169	75.630622	37.754884	9.504998
	>50K	36.046875	6102.859375	129.859375	11.156250
Unmarried	<=50K	39.948983	123.621874	32.652884	9.528843
	>50K	45.779343	5411.131455	150.173709	11.868545
Wife	<=50K	39.028090	282.980337	53.853933	9.567416
	>50K	40.789625	2980.096542	179.279539	11.475504
		hours.per.week			
relationship	income				
	Husband	<=50K	43.003096		
	>50K	46.434407			
Not-in-family	<=50K	40.380559			
	>50K	46.972053			
Other-relative	<=50K	37.135831			
	>50K	42.828571			
Own-child	<=50K	33.208996			
	>50K	43.281250			
Unmarried	<=50K	38.950984			
	>50K	46.201878			
Wife	<=50K	36.810393			
	>50K	38.466859			

For all of the relationship groups, the > 50k group has made a higher average capital gain amount.

```
[38]: print("Done by Brandon Cabrera")
race_pivot_table = pd.pivot_table(data= adult_census_data, index = ['race', 'income'], values= ['age', 'education.num', 'capital.gain', 'capital.loss', 'hours.per.week'])
```

```
race_pivot_table
```

Done by Brandon Cabrera

```
[38]:
```

race	income	age	capital.gain	capital.loss	\
Amer-Indian-Eskimo	<=50K	36.376984	203.535714	22.234127	
	>50K	39.411765	4045.647059	147.823529	
Asian-Pac-Islander	<=50K	36.151468	118.278207	48.755796	
	>50K	42.500000	4398.750000	203.677419	
Black	<=50K	36.989392	107.908201	40.806610	
	>50K	43.696721	3931.745902	159.806011	
Other	<=50K	32.780952	81.719048	46.338095	
	>50K	41.809524	10850.000000	89.857143	
White	<=50K	36.619723	155.210014	55.719860	
	>50K	44.055271	3899.515865	195.754642	

race	income	education.num	hours.per.week
Amer-Indian-Eskimo	<=50K	9.123016	39.781746
	>50K	11.088235	45.205882
Asian-Pac-Islander	<=50K	10.496136	38.925811
	>50K	12.407258	44.608871
Black	<=50K	9.306814	37.798450
	>50K	11.030055	44.540984
Other	<=50K	8.400000	39.633333
	>50K	11.523810	44.904762
White	<=50K	9.661307	39.553053
	>50K	11.611054	45.813715

The race with the highest mean age for the > 50k group, which despite having the most amount of people who are making > 50k. The races with fewer representation in the dataset have younger mean ages for their > 50k group.

```
[39]: print("Done by Brandon Cabrera")
sex_pivot_table = pd.pivot_table(data= adult_census_data, index = ['sex',  
                     'income'], values= ['age', 'education.num', 'capital.gain', 'capital.loss',  
                     'hours.per.week'])
sex_pivot_table
```

Done by Brandon Cabrera

```
[39]:
```

sex	income	age	capital.gain	capital.loss	education.num	\
Female	<=50K	36.231719	119.394348	46.251557	9.878777	
	>50K	41.964928	4084.820144	175.451439	11.812950	
Male	<=50K	36.841390	167.183352	57.909754	9.474328	
	>50K	44.305816	3912.098186	196.932145	11.570513	

```

hours.per.week
sex      income
Female   <=50K      36.423645
          >50K      40.897482
Male     <=50K      41.162042
          >50K      46.542683

```

The above pivot table allows me to answer the analysis question #9. For males who make > 50k, the average age is 44 years old, rounding to the nearest whole number. For females who make > 50k, the average age is 42 years old, rounding to the nearest whole number. On a side note, both males and females in the <= 50k group have the same average age of 36 without rounding, which is opposite of what we see with the > 50k group of females and males.

```
[40]: print("Done by Brandon Cabrera")
native_country_pivot_table = pd.pivot_table(data= adult_census_data, index = [
    ['native.country', 'income'], values= ['age', 'education.num', 'capital.gain',
    'capital.loss', 'hours.per.week']])
native_country_pivot_table
```

Done by Brandon Cabrera

```
[40]:
native.country  income      age  capital.gain  capital.loss  education.num \
Cambodia       <=50K      37.272727      543.545455      170.545455      7.727273
                  >50K      41.428571     1935.714286      0.000000     10.285714
Canada         <=50K      40.098592      70.605634      168.070423     10.239437
                  >50K      44.305556     4659.583333     105.250000     11.527778
China          <=50K      41.270833      119.041667      108.187500     10.354167
...
United-States  >50K      44.000429      3965.694353      194.689207     11.586276
Vietnam        <=50K      33.203390      252.406780      99.135593      9.508475
                  >50K      35.400000     5873.600000      0.000000     9.600000
Yugoslavia    <=50K      38.000000      0.000000      0.000000      9.700000
                  >50K      40.166667     926.000000      0.000000     10.333333

hours.per.week
native.country  income
Cambodia       <=50K      41.545455
                  >50K      40.000000
Canada         <=50K      38.704225
                  >50K      46.944444
China          <=50K      38.541667
...
United-States  >50K      45.750536
Vietnam        <=50K      38.152542
                  >50K      39.200000
Yugoslavia    <=50K      41.600000
                  >50K      49.500000
```

[80 rows x 5 columns]

Thailand's > 50k group has the highest average hours worked per week with 58 hours per week!

## 8 Testing Hypothesis

Now that we have answered our analysis questions let's do some hypothesis test for some of the observations we had for the questions to see if they are statistically significant and not by chance

Our first hypothesis test will help to confirm whether or not our observation for analysis question #1 is not by chance.

Let:

$\mu_0$ : means hour per week for the > 50k group

$\mu_1$ : mean hours per week for the  $\leq 50k$  group

Null Hypothesis: Both groups have the same average amount of hours worked per week,  $\mu_0 = \mu_1$

Alternate Hypothesis: The groups don't have the same average amount of hours worked per week,  $\mu_0 \neq \mu_1$

For this, we will use a Two-Sample Student T-Test:

- Test whether two independent samples are significantly different
- Assumptions:
  - Observations in each sample are independent and identically distributed (iid).
  - Observations in each sample are normally distributed.
  - Observations in each sample have the same variance.

```
[41]: from scipy.stats import ttest_ind
print("Done by Brandon Cabrera")
less_equal_to_50k_group =
    adult_census_data[adult_census_data["binary_income_over_50k"] == 0][["hours.
    per.week"]]
greater_than_50k_group =
    adult_census_data[adult_census_data["binary_income_over_50k"] == 1][["hours.
    per.week"]]
stat, p = ttest_ind(less_equal_to_50k_group, greater_than_50k_group)
if p > 0.05:
    print("Fail to reject null hypothesis")
else:
    print("Reject null hypothesis, there is evidence that the groups don't have
        the same average amount of hours worked per week")
```

Done by Brandon Cabrera

Reject null hypothesis, there is evidence that the groups don't have the same average amount of hours worked per week

We were able to reject the null hypothesis, meaning that the difference in the average hours worked per week for the two income groups is statistically significant. Our observation of the two groups having different means for the hours worked per week is likely to not be by chance.

Our second hypothesis test will be for the analysis question #9 is not by chance.

Let:

$\mu_0$ : mean age for males a part of the > 50k group

$\mu_1$ : mean age for females a part of the > 50k group

Null Hypothesis: Both groups are the same age,  $\mu_0 = \mu_1$

Alternate Hypothesis: The groups don't share the same age,  $\mu_0 \neq \mu_1$

For this, we will use a Two-Sample Student T-Test:

- Test whether two independent samples are significantly different
- Assumptions:
  - Observations in each sample are independent and identically distributed (iid).
  - Observations in each sample are normally distributed.
  - Observations in each sample have the same variance.

```
[42]: print("Done by Brandon Cabrera")
female_group = adult_census_data[(adult_census_data['binary_income_over_$50k'] == 1) & (adult_census_data['sex'] == 'Female')]['age']
male_group = adult_census_data[(adult_census_data['binary_income_over_$50k'] == 1) & (adult_census_data['sex'] == 'Male')]['age']
stat, p = ttest_ind(female_group, male_group)
if p > 0.05:
    print("Fail to reject null hypothesis")
else:
    print("Reject null hypothesis, there is evidence that the groups don't share the same mean age")
```

Done by Brandon Cabrera

Reject null hypothesis, there is evidence that the groups don't share the same mean age

We were able to reject the null hypothesis. This means that our observation for analysis question #9 is in line with the hypothesis test since we found that the average age for women a part of the > 50k group was different than the average age of males a part of the > 50k group.

## 9 Conclusions

### 9.0.1 Analysis question #1

Q: What was the average number of hours worked for a person who made > 50k, and what was the average number of hours worked for a person who made <= to 50k?

A: The average number of hours worked for a person who made  $> 50k$  is 45.7 hours worked per week. The average number of hours worked for a person who made  $\leq 50k$  is 39.3 hours worked per week.

### 9.0.2 Analysis question #2

Q: For a person whose highest education level is only high school, who makes  $> 50k$ , if any, what is the average number of hours worked?

A: For a person whose highest education level is HS-grad and is a part of the  $> 50k$  group, the average age is 45 years old, rounding to the nearest whole number.

### 9.0.3 Analysis question #3

Q: What is the average education level for a person making  $> 50k$  and what is the average education level for a person making  $\leq 50k$ ?

A: The average education number for a person making  $> 50k$  is 12 if we round up to the nearest number, which is equivalent to Assoc-acdm in terms of education level. The average education number for a person making  $\leq 50k$  is 10, meaning they have some college.

### 9.0.4 Analysis question #4

Q: What is the correlation between capital gain and whether or not a person makes over  $50k$ ?

A: The correlation between capital gain and a person making over  $> 50k$  is 0.221, indicating a weak positive relationship. This makes sense since not everyone who is making money from investments might be making a profit compared to their losses, or the money made isn't much. The correlation between capital gain and a person making  $\leq 50k$  is -0.221 indicating, a weak negative relationship.

### 9.0.5 Analysis question #5

Q: What marital status has the most people making over  $50k$ , and what marital status has the most people making less than or equal to  $50k$ ?

A: The marital status with the most amount of people making  $> 50k$  is married-civ-spouse. The marital status with the most amount of people making  $\leq 50k$  is Never-married.

### 9.0.6 Analysis question #6

Q: What numerical column has the highest amount of correlation, in terms of magnitude, with whether or not a person made  $> 50k$  and whether they make  $\leq 50k$ ?

A: The column with the highest amount of correlation, in terms of magnitude, for both binary columns is education.num with a value of 0.335. This is a weak relationship, indicating that as a person's education level becomes higher, they are more likely to make  $> 50k$ . If we look at it from the perspective of making  $\leq 50k$ , -0.335 indicates it indicates that as a person's education level becomes higher than they are less likely to make  $\leq 50k$ .

### **9.0.7 Analysis question #7**

Q: What workclass has the most people making  $\leq 50k$ , and what work class has the most people making  $> 50k$ ?

A: The private workclass has the most amount of people making  $> 50k$ , as well as the most amount of people making less  $\leq 50k$ .

### **9.0.8 Analysis question #8**

Q: Do women or men tend to make  $> 50k$  more than the other, and do women or men tend to make  $\leq 50k$ ?

A: The sex with the most amount of people who make  $> 50k$  and the most of people who make  $\leq 50k$  is the male sex.

### **9.0.9 Analysis question #9**

Q: What is the average age of men who make  $> 50k$ , and what is the average age of women who make  $> 50k$ ?

A: For males who make  $> 50k$ , the average age is 44 years old, rounding to the nearest whole number. For females who make  $> 50k$ , the average age is 42 years old, rounding to the nearest whole number.

### **9.0.10 Analysis question #10**

Q: Which race tends to make  $> 50k$ , more than the other races, and what race tends to make  $\leq 50k$  more than the other races?

A: The race with the most amount of people making  $> 50k$  is the white race, and it's also the race with the most amount of people making  $\leq 50k$ . It's worth noting that I am considering only the frequency and not the percentages, which would be a fairer assessment of which race tends to make more.